

MULTI-LLM BENCHMARKING & EVALUATION PLATFORM

Technical Documentation & API Reference Manual

1. SYSTEM OVERVIEW & ARCHITECTURE

The Multi-LLM Benchmarking & Evaluation Platform is an enterprise-grade Django REST API application designed to dynamically process user prompts, evaluate AI model outputs through a double-blind jury system, and resolve a definitive winner utilizing a statistical consensus-scoring mechanism.

Core Architectural Pipeline:

1. **Ingress Validation:** The client submits a payload containing a prompt. The platform filters and validates the input length via serialization constraints.
2. **Asynchronous Orchestration:** Using Python's `asyncio` engine, the backend dispatches non-blocking concurrent API requests to multiple upstream LLM providers (OpenAI, Gemini, DeepSeek, Groq).
3. **Double-Blind Anonymization:** Model identities are stripped dynamically and mapped to uppercase token labels (`Response A`, `Response B`, etc.) to guarantee unbiased downstream evaluation.
4. **Independent Evaluation Jury:** Anonymized outputs are parsed through multi-agent AI Judges (Gemini and Groq) evaluating core linguistic vectors: Accuracy, Relevance, Completeness, Clarity, and Conciseness.
5. **Statistical Consensus Blending :** Text structures undergo TF-IDF vectorization and Cosine Similarity checking to isolate aberrant hallucinations or outliers, updating model scores mathematically before crowning the session winner.
6. **Relational Persistence:** Complete session audit logs are atomically saved inside a relational storage engine (PostgreSQL) using optimized query techniques.

2. ENVIRONMENT CONFIGURATION (`.env`)

To run this application locally or in production, create a `.env` file in the root project directory. Copy the sample block below and replace the placeholders with your actual API gateway credentials.

```
# Server Deployment Configuration
```

```
DEBUG=True
```

```
SECRET_KEY=django-insecure-mock-key-evaluation-platform-2026
```

Upstream LLM Provider API Keys

OPENROUTERAPI_KEY=sk-proj-mockOpenAiKeyForTestingPurposesOnly12345

GEMINI_API_KEY=AlzaSyMockGeminiKeyGenAIPlatformRegistration2026

GROQ_API_KEY=gsk_mockGroqHighPerformanceInferenceEngineKey99

Database Configuration (Uncomment for PostgreSQL Setup)

DB_NAME=multi_llm_bench_db

DB_USER=postgres

DB_PASSWORD=your_secure_db_password

DB_HOST=localhost

DB_PORT=5432\

3. API SPECIFICATIONS & DOCUMENTATION

ENDPOINT 1: Submit Prompt & Execute Evaluation Benchmark

- **URL:** `/api/evaluate/`
- **Method:** `POST`
- **Headers:** `Content-Type: application/json`

http://127.0.0.1:8000/api/evaluate/

Server response

Code	Details
200	Response body

Undocumented

```
{
  "session_id": 1,
  "prompt": "what is blackhole",
  "winner": "Openrouter-gpt-oss-20b:free",
  "final_score": 9.5,
  "all_responses": [
    {
      "label": "A",
      "model_name": "Groq-llama-3.1-8b-instant",
      "response_text": "A black hole is a region in space where the gravitational pull is so strong that nothing, including light, can escape. It is formed when a massive star collapses in on itself and its gravity becomes so strong that it warps the fabric of spacetime around it.\n\nHere's a simplified explanation of how a black hole is formed:\n\n1. **Massive Star Collapse**: A massive star runs out of fuel and dies. Its core collapses under its own gravity, causing a massive amount of matter to be compressed into an incredibly small space.\n\n2. **Singularity Formation**: The collapse creates a point of infinite density and zero volume, known as a singularity. This is the center of the black hole.\n\n3. **Event Horizon Formation**: As matter falls towards the singularity, it creates a boundary called the event horizon. Once something crosses the event horizon, it is trapped by the black hole's gravity and cannot escape.\n\n4. **Gravitational Pull**: The event horizon marks the point of no return. Any matter or radiation that enters the event horizon is pulled towards the singularity, never to escape.\n\n**Characteristics of Black Holes**:\n\n1. **Mass**: Black holes have mass, which determines the strength of their gravity.\n\n2. **Spin**: Black holes can rotate, which affects the direction of their jets of energy and the way they interact with their surroundings.\n\n3. **Charge**: Black holes can have an electric charge, which affects their behavior in the presence of other charged objects.\n\n**Types of Black Holes**:\n\n1. **Stellar Black Holes**: Formed from the collapse of individual stars, these black holes have masses up to about 20 times that of the sun.\n\n2. **Supermassive Black Holes**: Found at the centers of galaxies, these black holes have masses millions or even billions of times that of the sun.\n\n3. **Intermediate-Mass Black Holes**: Black holes with masses between those of stellar and supermassive black holes.\n\n**Observational Evidence for Black Holes**:\n\n1. **X-rays and Gamma Rays**: Telescopes can detect X-rays and gamma rays emitted by hot gas swirling around black holes.\n\n2. **Radio Waves**: Radio telescopes can detect radio waves emitted by matter as it spirals into a black hole.\n\n3. **Gravitational Waves**: The detection of gravitational waves by LIGO and VIRGO provide strong evidence for the existence of black holes.\n\n**Interesting Facts about Black Holes**:\n\n1. **Information Paradox**: The laws of quantum mechanics suggest that information cannot be destroyed, but the laws of general relativity suggest that information is lost in a black hole. This paradox is still an open question in physics.\n\n2. **Black Hole Entropy**: Black holes have a temperature and entropy, which is related to the area of their event horizon.\n\n3. **Black Hole Complementarity**: The information paradox is resolved by considering that information that falls into a black hole is not lost, but is instead stored on the event horizon and can be retrieved in some way."
    }
  ]
}
```

Response headers

Without addition of consensus

POST http://127.0.0.1:8000/api/evaluate/ Send

Docs Params Authorization Headers 9 Body Scripts Tests Settings Cookies

Headers 8 hidden

Key	Value	Description	Bulk Edit	Presets	...
Content-Type	application/json				

Body Cookies Headers 10 Test Results 200 OK • 15.41 s • 5.05 KB

JSON Preview Visualize

```
1 {
2   "session_id": 2,
3   "prompt": "briefly tell what is machine learning?",
4   "winner": "Groq-llama-3.1-8b-instant",
5   "final_score": 9.0,
6   "all_responses": [
7     {
8       "label": "A",
9       "model_name": "Groq-llama-3.1-8b-instant",
10      "response_text": "**Machine Learning (ML)** is a subset of artificial intelligence (AI) that involves training algorithms",
11      "response_time_ms": 1085.81,
12      "groq_score": 9,
13      "gemini_score": 9.0,
14      "final_score": 9.0
15    },
16    {
17      "label": "B",
18      "model_name": "Groq-qwen/qwen3-32b",
```

With consensus

POST http://127.0.0.1:8000/api/evaluate/ Send

Docs Params Authorization Headers 9 Body Scripts Tests Settings Cookies

none form-data x-www-form-urlencoded raw binary GraphQL JSON Schema Beautify

Body Cookies Headers 10 Test Results 200 OK • 34.13 s • 10.55 KB

JSON Preview Visualize

```
1 {
2   "session_id": 6,
3   "prompt": "AI KYA HA?",
4   "winner": "OpenRouter-gpt-oss-20b:free",
5   "final_score": 6.98,
6   "judges_used": {
7     "groq": true,
8     "gemini": true
9   },
10  "consensus": {
11    "scores": {
12      "A": 0.0954,
13      "B": 0.1083,
14      "C": 0.1022,
15      "D": 0.0916
16    },
17    "outliers": [],
18    "confidence": "low",
```

ENDPOINT 2: Historical Evaluation Sessions Log

- URL: </api/history/>
- Method: GET
- Headers: Accept: application/json

- **Description:** Retrieves a registry containing all historical model comparison evaluation records. The ORM query layers execute structural `prefetch_related` wrappers to eliminate database performance bottlenecks (N+1 query loops).

The screenshot shows a web browser with a REST client interface. The curl command is as follows:

```
curl -X 'GET' \
  'http://127.0.0.1:8000/api/history/' \
  -H 'accept: application/json' \
  -H 'X-CSRFToken: 163GPMqxkVhYpUCJnQG00YdFSBYzmGf1XjELae5mQKz0rgSExf36bTEbIisn1'
```

The Request URL is `http://127.0.0.1:8000/api/history/`.

The Server response is 200 OK. The Response body is a JSON array with one object:

```
{
  "id": 1,
  "prompt": "what is blackhole",
  "winner_model": "OpenRouter-gpt-oss-20b-free",
  "winner_score": 9.5,
  "responses": [
    {
      "id": 1,
      "model_name": "Groq-llama-3.1-8b-instant",
      "anonymous_label": "A",
      "response_text": "A black hole is a region in space where the gravitational pull is so strong that nothing, including light, can escape. It is formed when a massive star collapses in on itself and its gravity becomes so strong that it warps the fabric of spacetime around it. Here's a simplified explanation of how a black hole is formed:\n\n1. **Massive Star Collapse**: A massive star runs out of fuel and dies. Its core collapses under its own gravity, causing a massive amount of matter to be compressed into an incredibly small space.\n\n2. **Singularity Formation**: The collapse creates a point of infinite density and zero volume, known as a singularity. This is the center of the black hole.\n\n3. **Event Horizon Formation**: As matter falls towards the singularity, it creates a boundary called the event horizon. Once something crosses the event horizon, it is trapped by the black hole's gravity and cannot escape.\n\n4. **Gravitational Pull**: The event horizon marks the point of no return. Any matter or radiation that enters the event horizon is pulled towards the singularity, never to escape.\n\n**Characteristics of Black Holes**:\n\n1. **Mass**: Black holes have mass, which determines the strength of their gravity.\n\n2. **Spin**: Black holes can rotate, which affects the structure of their event horizon and the paths of objects nearby."
    }
  ]
}
```

The screenshot shows a REST client interface with a GET request to `http://127.0.0.1:8000/api/history/`. The response is 200 OK, 118 ms, 17.26 KB. The response body is a JSON array with one object:

```
{
  "id": 2,
  "prompt": "briefly tell what is machine learning?",
  "winner_model": "Groq-llama-3.1-8b-instant",
  "winner_score": 9.0,
  "responses": [
    {
      "id": 5,
      "model_name": "Groq-llama-3.1-8b-instant",
      "anonymous_label": "A",
      "response_text": "**Machine Learning (ML)** is a subset of artificial intelligence (AI) that involves training algorithms to learn from data and make predictions or decisions based on the learned patterns.",
      "response_time_ms": 1085.8101844787598,
      "final_score": 9.0,
      "judge_scores": [
        {
          "judge_name": "Groq-Judge",
          "score": 9.0
        },
        {
          "judge_name": "Gemini-Judge",
          "score": 9.0
        }
      ]
    }
  ]
}
```

ENDPOINT 3: Isolated Session Details

- **URL:** `/api/history/<id>/` (e.g., `/api/history/1/`)
- **Method:** GET
- **Headers:** Accept: application/json

- **Description:** Pulls isolated deep audit logs for a single target execution session parsed via its integer primary key ID. Validates route data types, rendering automatic **404 Not Found** messages if indices do not exist.

```

Curl
curl -X 'GET' \
  'http://127.0.0.1:8000/api/history/' \
  -H 'accept: application/json' \
  -H 'X-CSRFToken: 1636PMqxkMwVypUCJnQG00YdFSBYzmgf1XjELae5mQKz0rg5Exf36bTeXbI1sn1'

```

Request URL

http://127.0.0.1:8000/api/history/

Server response

Code Details

200

Response body

```

{
  "id": 1,
  "prompt": "what is blackhole",
  "winner_model": "OpenRouter-gpt-oss-20b-free",
  "winner_score": 9.5,
  "responses": [
    {
      "id": 1,
      "model_name": "Groq-llama-3.1-8b-instant",
      "anonymous_label": "A",
      "response_text": "A black hole is a region in space where the gravitational pull is so strong that nothing, including light, can escape. It is formed when a massive star collapses in on itself and its gravity becomes so strong that it warps the fabric of spacetime around it. Here's a simplified explanation of how a black hole is formed:\n\n1. **Massive Star collapse**: A massive star runs out of fuel and dies. Its core collapses under its own gravity, causing a massive amount of matter to be compressed into an incredibly small space.\n\n2. **Singularity Formation**: The collapse creates a point of infinite density and zero volume, known as a singularity. This is the center of the black hole.\n\n3. **Event Horizon Formation**: As matter falls towards the singularity, it creates a boundary called the event horizon. Once something crosses the event horizon, it is trapped by the black hole's gravity and cannot escape.\n\n4. **Gravitational Pull**: The event horizon marks the point of no return. Any matter or radiation that enters the event horizon is pulled towards the singularity, never to escape.\n\n**Characteristics of Black Holes**:\n\n1. **Mass**: Black holes have mass, which determines the strength of their gravity.\n\n2. **Spin**: Black holes can rotate, which affects the distribution of their spacetime curvature.\n\n3. **Charge**: Black holes can have an electric charge, which affects their interaction with other charged particles."
    }
  ]
}

```

GET ▼ http://127.0.0.1:8000/api/history/2 Send ▼

Docs Params Authorization Headers 9 Body ● Scripts Tests Settings Cookies

Body Cookies Headers 10 Test Results ↺ 200 OK • 93 ms • 5.63 KB • 🌐 • ⋮

{} JSON ▼ ▶ Preview 🖨 Visualize ▼ ↺ ≡ 🔍 📄 🔗

```

1 {
2   "id": 2,
3   "prompt": "briefly tell what is machine learning?",
4   "winner_model": "Groq-llama-3.1-8b-instant",
5   "winner_score": 9.0,
6   "responses": [
7     {
8       "id": 5,
9       "model_name": "Groq-llama-3.1-8b-instant",
10      "anonymous_label": "A",
11      "response_text": "**Machine Learning (ML)** is a subset of artificial intelligence (AI) that involves training algorithms
12      "response_time_ms": 1085.8101844787598,
13      "final_score": 9.0,
14      "judge_scores": [
15        {
16          "judge_name": "Groq-Judge",
17          "score": 9.0
18        },
19        {
20          "judge_name": "Gemini-Judge",
21          "score": 9.0
22        }
23      ]
24    }
25  ]
26 }

```

ENDPOINT 4: Global Performance Leaderboard Rankings

- **URL:** /api/leaderboard/
- **Method:** GET
- **Headers:** Accept: application/json

- Description:** Compiles real-time analytical metrics by grouping historical evaluation blocks at the database layer. Leverages Django **Count** and **Avg** functions to compute global model performance stats, ordering the output dynamically using total match wins.

Request URL: `http://127.0.0.1:8000/api/leaderboard/`

Server response: 200

Response body:

```
[
  {
    "model_name": "OpenRouter-gpt-oss-20b:free",
    "total_wins": 1,
    "average_score": 9.5,
    "total_evaluations": 1
  },
  {
    "model_name": "Groq-llama-3.1-8b-instant",
    "total_wins": 0,
    "average_score": 9,
    "total_evaluations": 1
  },
  {
    "model_name": "OpenRouter-nex-n2-pro:free",
    "total_wins": 0,
    "average_score": 7,
    "total_evaluations": 1
  },
  {
    "model_name": "Groq-qwen/qwen3-32b",
    "total_wins": 0,
    "average_score": 4,
    "total_evaluations": 1
  }
]
```

Response headers:

GET `http://127.0.0.1:8000/api/leaderboard/` Send

Docs Params Authorization Headers 9 Body Scripts Tests Settings Cookies

Body Cookies Headers 10 Test Results 200 OK • 131 ms • 724 B • ...

{ } JSON Preview Visualize

```
1 [
2   {
3     "model_name": "Groq-llama-3.1-8b-instant",
4     "total_wins": 1,
5     "average_score": 9.0,
6     "total_evaluations": 2
7   },
8   {
9     "model_name": "OpenRouter-gpt-oss-20b:free",
10    "total_wins": 1,
11    "average_score": 8.0,
12    "total_evaluations": 2
13  },
14  {
15    "model_name": "OpenRouter-nex-n2-pro:free",
16    "total_wins": 0,
17    "average_score": 7.25,
18    "total_evaluations": 2
19  },
20  {
21    "model_name": "Groq-qwen/qwen3-32b",
22    "total_wins": 0,
```

al 0 0 0 Globals Vault Tools

4. ADVANCED FEATURE OVERVIEW: STATISTICAL CONSENSUS ENGINE

To meet the evaluation **Bonus Challenge** criteria, the platform implements an embedded mathematical tracking ecosystem inside **consensus.py** instead of relying blindly on simple judge score averages.

- **Semantic Similarity Vectoring:** Generated model response texts are parsed dynamically using TF-IDF (Term Frequency-Inverse Document Frequency) matrices. The system then calculates cross-token agreement distributions using Cosine Similarity checks.
- **Algorithmic Outlier Isolation:** Any model response displaying extreme deviations from structural similarity averages (Standard Deviation boundary factors exceeding $1.5 \times \sigma$ combined with raw similarity margins falling below \$0.15\$) is flagged as an aberrant outlier or hallucination.
- **Weighted Score Penalty Mitigation:** When an outlier node is isolated, its judge criteria calculations are dynamically down-weighted to minimize its chance of being selected as the winner, preventing corrupted tokens from gaming the leaderboard rankings.

$$\text{Adjusted Score} = (\text{Raw Judge Score} \times (1 - w)) + (\text{Consensus Scale} \times w)$$

```
Adjusted Score = (Raw Judge Score x (1 - w)) + (Consensus Scale x w)
```

```
Where: w = 0.20 (consensus weight = 20%)
```

```
1 - w = 0.80 (judge weight = 80%)
```

```
Example:
```

```
Judge Score = 8.5
```

```
Consensus = 0.91 --> Scaled = 0.91 x 10 = 9.1
```

```
Adjusted = (8.5 x 0.80) + (9.1 x 0.20)
```

```
= 6.80 + 1.82
```

```
= 8.62 <-- Final Score
```

Challenges Faced and Resolutions

1. OpenAI API Quota Exceeded

Issue: You exceeded your current quota, please check your plan and billing details.

Cause: The API key was valid, but the OpenAI account had no available quota or active billing.

Resolution: OpenAI integration done using openrouter api.

```
POST http://127.0.0.1:8000/api/submit/ 200 OK • 1.44 s • 1.18 KB • ...
Body Cookies Headers 10 Test Results
{ } JSON Preview Visualize
3  "responses": [
4  {
5    "model_name": "openai",
6    "response": "API Error: You exceeded your current quota, please check your plan and billing details. For more information v
7    "response_time": 1.39
8  },
9  {
10   "model_name": "gemini",
11   "response": "API Error: models/gemini-pro is not found for API version v1beta, or is not supported for generateContent. Ca
12   "response_time": 0.6
13 },
14 {
15   "model_name": "deepseek",
16   "response": "API Error: Insufficient Balance",
17   "response_time": 0.78
18 },
19 {
20   "model_name": "groq",
21   "response": "API Error: The model 'llama3-70b-8192' has been decommissioned and is no longer supported. Please refer to ht
22   "response_time": 0.28
23 },
24 ]
```

2. Gemini Model Not Found

Issue:models/gemini-1.5-flash is not found for API version v1b

Cause:The application was using an outdated Gemini model that was no longer available for the provided API version.

Resolution:Available models were retrieved using the Gemini Models API and the integration was updated to use gemini-2.5-flash

3. Gemini Free Tier Quota Exceeded

Issue:Quota exceeded for metric generate_content_free_tier_requests

Cause:the Gemini free-tier request limit had been reached.

Resolution:Retry handling was considered and alternative judge models were used when quota limits were encountered.

4. DeepSeek Insufficient Balance

POST

https://api.deepseek.com/chat/completions

Docs

Params

Authorization

Headers 9

Body

Scripts

Tests

Se

Query Params

Key	Value
Key	Value

Body

Cookies

Headers 14

Test Results

{ } JSON

Preview

Debug with AI

```
1 {
2   "error": {
3     "message": "Insufficient Balance",
4     "type": "unknown_error",
5     "param": null,
6     "code": "invalid_request_error"
7   }
8 }
```

https://api.deepseek.com/chat/completions

Save

Share

GET

https://api.deepseek.com/chat/completions

Send

Docs

Params

Authorization

Headers 9

Body

Scripts

Tests

Settings

Cookies

none

form-data

x-www-form-urlencoded

raw

binary

GraphQL

JSON

Schema

Beautify

```
6   "content": "Explain Django ORM in simple words"
7   }
8 ]
9 }
```

Body

Cookies

Headers 9

Test Results

403 Forbidden • 1.08 s • 1.23 KB

HTML

Preview

Try with different auth credentials

```
1 <!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN" "http://www.w3.org/TR/html4/loose.dtd">
2 <HTML>
3
4 <HEAD>
5   <META HTTP-EQUIV="Content-Type" CONTENT="text/html; charset=iso-8859-1">
6   <TITLE>ERROR: The request could not be satisfied</TITLE>
7 </HEAD>
8
9 <BODY>
10  <H1>403 ERROR</H1>
11  <H2>The request could not be satisfied.</H2>
12  <HR noshade size="1px">
13  Bad request.
14  We can't connect to the server for this app or website at this time. There might be too much traffic or a
15  configuration error. Try again later, or contact the app or website owner.
16  <BR clear="all">
```

0 0 0

Globals Vault Tools

Issue:Insufficient Balance

Cause:The DeepSeek account did not have sufficient credits to process requests.

Resolution:DeepSeek was removed from active evaluation and replaced with Groq-hosted models.

5. Groq Model Deprecation

Issue:The model llama-3.1-70b-versatile has been decommissioned

Cause:The selected Groq model had been officially deprecated.

Resolution:The deprecated model was replaced with supported models:

llama-3.1-8b-instant

qwen/qwen3-32b

openai/gpt-oss-20b


6. Django Async ORM Error

Issue:SynchronousOnlyOperation:You cannot call this from an async context

Cause:Django ORM queries were executed directly inside asynchronous functions.

Resolution:Database operations were moved outside async contexts or wrapped using asynchronous-safe techniques.

7.Swagger UI Schema Mapping Issue



Error: Swagger UI displayed *"No parameters"* for the `/api/evaluate/` POST endpoint.

Cause: The endpoint was powered by a standard Django REST Framework `APIView`. By default, `drf-yasg` (Swagger generator) cannot auto-detect the expected input data payload structure for raw `APIView` classes unless explicitly documented. Because there was no explicit link between the view and the data validator (`PromptInputSerializer`), Swagger rendered an empty parameter block.

Resolution: Decorated the view function with `@swagger_auto_schema(request_body=PromptInputSerializer)`. This explicitly mapped the serializer schema to the Swagger engine, turning the static documentation into an interactive JSON input portal with an "Edit Value" field.

127.0.0.1:8000/swagger/

POST /evaluate/ evaluate_create

POST /api/evaluate/

Parameters

Cancel

Name	Description
data * required	
object	Edit Value Model
(body)	

PromptInput {

prompt*

Prompt { string

title: Prompt

maxLength: 2000

minLength: 5

}

Execute

Responses

Response content type application/json

8.hugging face api did not worked

https://api-inference.huggingface.co/models/meta-llama/Llama-3.2-3B-Instruct/v1/chat/completions Save Share

POST https://api-inference.huggingface.co/models/meta-llama/Llama-3.2-3B-Instruct/v1/chat/completions Send

Docs Params Authorization Headers 9 Body Scripts Tests Settings Cookies

none form-data x-www-form-urlencoded raw binary GraphQL JSON Schema Beautify

```

1 {
2   "messages": [
3     {
4       "role": "user",
5       "content": "Explain Django ORM in simple words"
6     }
7   ]
8 }

```

Response

Could not send request


Error: getaddrinfo ENOTFOUND api-inference.huggingface.co Debug with AI

View in console • Learn about troubleshooting requests

al 0 0 0 Globals Vault Tools

getaddrinfo ENOTFOUND api-inference.huggingface.co — Could not send request.

9.aiml api did not work



404 Not Found — Model not found. NotFoundException at /v1/chat/completions.

Final Integrated Models

LLM Models (Response Generation)

Model	Provider	API	Label
llama-3.1-8b-instant	Groq	api.groq.com	A
qwen/qwen3-32b	Groq	api.groq.com	B
nex-agi/nex-n2-pro:free	OpenRouter	openrouter.ai	C
gpt-oss-20b:free	OpenRouter	openrouter.ai	D

Judge Models (Evaluation)

Judge	Model Used	Scoring Criteria
Groq Judge	llama-3.1-8b-instant	Accuracy, Relevance, Completeness, Clarity, Conciseness

Gemini Judge	gemini-2.5-flash	Accuracy, Relevance, Completeness, Clarity, Conciseness
---------------------	------------------	---

Both judges independently evaluate anonymized model responses (labeled A, B, C, D). Final score = Average(Groq Judge Score, Gemini Judge Score). Winner = highest adjusted score after consensus blending.